
Natural Language to SQL in Healthcare: Bridging Analytics Maturity Gaps, Workforce Turnover, and Technical Barriers Through Conversational AI Platforms

Samuel T Harrold, Yuimedi

December 2025

This research examines the evidence for implementing conversational AI platforms in healthcare analytics, addressing three critical challenges: low healthcare analytics maturity, workforce turnover with institutional memory loss, and technical barriers in natural language to SQL generation. Through review of peer-reviewed benchmarking studies and industry implementations, we demonstrate that natural language interfaces can democratize analytics access while preserving institutional knowledge. Healthcare-specific text-to-SQL benchmarks show significant progress, though current models are “not yet sufficiently accurate for unsupervised use” in clinical settings. Healthcare IT staff turnover of ~34%—the highest among IT sectors—creates institutional memory loss, while low-code implementations demonstrate significant efficiency gains and cost savings. The convergence of technical advances in NL2SQL generation, analytics maturity challenges in healthcare organizations, and workforce turnover creates both urgent need and strategic opportunity for conversational AI platforms with appropriate governance. This paper contributes a three-pillar analytical framework and positions healthcare conversational AI as a knowledge portal architecture for institutional memory preservation.

Contents

1 Executive Summary	4
2 Introduction	5
2.1 Background	5
2.2 Problem Statement	5
2.2.1 Low Healthcare Analytics Maturity	6
2.2.2 Technical Barriers to Data Access	6
2.2.3 Institutional Memory Loss from Workforce Turnover	6
2.3 Objectives	7
2.3.1 Primary Objective	7
2.3.2 Secondary Objectives	7
2.3.3 Non-Goals	7
2.4 Contributions	8
2.5 Document Structure	8
3 Methodology	8
3.1 Review Approach	8
3.2 Literature Search	9
3.3 Source Selection	10
3.4 Evidence Synthesis	10
3.5 Methodological Limitations	10
4 Literature Review: Natural Language Analytics in Healthcare - Evidence for Institutional Memory Preservation	11
4.1 1. Current State of Natural Language to SQL Generation	11
4.1.1 Evolution and Technical Advances	11

4.1.2 Healthcare-Specific Challenges	12
4.1.3 Promising Approaches and Limitations	12
4.2 2. State of Healthcare Analytics Maturity	13
4.2.1 Low Organizational Maturity	13
4.2.2 Barriers to Analytics Adoption	13
4.2.3 The Analytics Skills Gap	14
4.3 3. Healthcare Workforce Turnover and Knowledge Loss	14
4.3.1 Turnover Rates and Financial Impact	14
4.3.2 Institutional Memory Loss	15
4.3.3 Traditional Approaches Inadequate	15
4.4 4. Integration of Evidence: The Case for Conversational AI	15
4.4.1 Bridging Technical and Domain Expertise	15
4.4.2 Knowledge Preservation Mechanisms	16
4.4.3 Empirical Support for Low-Code Healthcare Solutions	16
4.5 5. Implications for Healthcare Organizations	17
4.5.1 Strategic Alignment with Industry Trends	17
4.5.2 Return on Investment Evidence	17
4.5.3 Risk Mitigation Through Knowledge Preservation	17
4.6 6. Gaps in Current Literature	18
4.7 7. Why the Problem Persists	18
4.7.1 Failed Standardization Approaches	18
4.7.2 Structural Disincentives in the Technology Market	19
4.7.3 Deployment Constraint Mismatch	19
5 Proposed Solution: Conversational AI Platforms for Healthcare Analytics	19
5.1 Solution Overview	20
5.2 Core Capabilities	20
5.2.1 Healthcare-Optimized Natural Language Processing	20
5.2.2 Healthcare Knowledge Portal	22
5.2.3 Progressive Analytics Maturity Development	22
5.2.4 Adaptive Query Generation and Optimization	23
5.3 Implementation Framework	23
5.3.1 Foundation and Integration (Months 1-3)	23
5.3.2 Knowledge Capture and Learning (Months 4-6)	24
5.3.3 Democratization and Scale (Months 7-12)	24
5.4 Risk Mitigation and Quality Assurance	25
5.4.1 Data Quality and Accuracy	25
5.4.2 Change Management and Adoption	25
5.4.3 Regulatory Compliance and Security	25
6 Evaluation: Empirical Evidence from Healthcare Implementations	26
6.1 Academic Study Results	26
6.1.1 LLM Benchmarking in Healthcare	26
6.1.2 Healthcare Text-to-SQL Benchmarks	26
6.1.3 RAG for Healthcare Queries	26
6.1.4 NLP in Healthcare	27

6.2 Real-World Case Studies	27
6.2.1 Berkshire Healthcare NHS Trust	27
6.3 Economic Impact Analysis	27
6.3.1 Return on Investment Evidence	27
6.3.2 Market Validation and Growth	28
7 Discussion	28
7.1 Strengths of the Evidence Base	28
7.1.1 Validated Benchmarking Data	28
7.1.2 Real-World Implementation Evidence	29
7.1.3 Addresses Multiple Challenges Simultaneously	29
7.1.4 Strong Economic Justification	29
7.1.5 Honest Assessment of Limitations	29
7.2 Limitations and Constraints	29
7.2.1 Implementation Complexity	29
7.2.2 Context-Specific Customization Requirements	30
7.2.3 Long-Term Outcome Uncertainties	30
7.2.4 Governance and Quality Assurance Challenges	30
7.2.5 Specialty-Specific Application Gaps	30
7.2.6 Methodological Considerations	30
7.3 Future Research Directions	31
7.3.1 Short-Term Research Priorities (<1 year)	31
7.3.2 Medium-Term Research Priorities (1-2 years)	31
7.3.3 Long-Term Research Priorities (>2 years)	31
7.4 Implications for Healthcare Organizations	31
7.4.1 Strategic Imperative	32
7.4.2 Implementation Approach	32
7.4.3 Competitive Advantage	32
8 Conclusion	32
8.1 The Dual Dimensions of Harm	32
8.2 Summary of Contributions	33
8.3 Key Findings	34
8.4 Implications for Organizational Assessment	34
8.5 Future Research Directions	35
8.6 Closing Reflection	35
9 Acknowledgments	36
10 Author Contributions	36
11 Competing Interests	36
12 Data Availability	36
13 Code Availability	36
14 Funding	37

15 References	37
15.1 Academic Sources	37
15.2 Industry Sources	40
16 Appendices	41
16.1 Appendix A: Healthcare Analytics Glossary	41
16.2 Appendix B: HIMSS Analytics Maturity Assessment Model (AMAM) Stages	42
16.3 Appendix C: Healthcare NL2SQL Query Examples	43
16.3.1 Patient Population Analysis	43
16.3.2 Operational Metrics	43
16.3.3 Quality Metrics	43

1 Executive Summary

Healthcare organizations face a critical convergence of challenges that threaten their ability to leverage data for improved patient outcomes and operational efficiency. This research examines evidence supporting conversational AI platforms as a strategic solution to three interconnected problems: persistently low healthcare analytics maturity, devastating institutional memory loss from workforce turnover, and technical barriers preventing clinical professionals from accessing their own data.

Through systematic review of academic and industry sources, we demonstrate that few healthcare organizations worldwide have achieved advanced analytics maturity, while healthcare IT staff turnover of 34% [A10]—the highest among IT sectors—creates institutional memory loss with knowledge loss costs reaching three times annual salary budgets [A24]. Simultaneously, natural language to SQL (NL2SQL) technologies have matured sufficiently to address healthcare’s unique technical barriers, though current models are “not yet sufficiently accurate for unsupervised use” in clinical settings [A6].

Conversational AI platforms directly address this convergence by democratizing analytics access through natural language interfaces while functioning as healthcare knowledge portals [A25] that preserve institutional knowledge through encoded expertise. Evidence from healthcare implementations shows significant improvements in efficiency, with organizations like Berkshire Healthcare NHS Trust reporting over 800 citizen developers creating solutions [I4], and Forrester Research documenting 206% ROI from low-code implementations [I5].

The strategic imperative is clear: healthcare organizations must adopt conversational AI platforms to preserve institutional memory, advance analytics maturity, and enable evidence-based decision making in an era of unprecedented workforce challenges.

2 Introduction

2.1 Background

Healthcare analytics has emerged as a critical capability for improving patient outcomes, reducing costs, and enhancing operational efficiency. While healthcare organizations must balance cost management, regulatory compliance, and operational efficiency, these concerns serve a primary institutional imperative: delivering high-quality patient care. Analytics initiatives that fail to advance this core mission—or worse, that divert resources and attention without improving care delivery—represent a misalignment with healthcare’s fundamental purpose.

However, the sector faces unique challenges that distinguish it from other data-intensive industries. Unlike technology or financial services, healthcare combines complex clinical workflows, extensive regulatory requirements, and a workforce with limited technical training but deep domain expertise.

The Healthcare Information Management Systems Society (HIMSS) Analytics Maturity Assessment Model (AMAM) provides the industry standard for measuring healthcare analytics capabilities across seven stages, from basic data collection to advanced predictive modeling and AI integration. Recent assessments reveal a sobering reality: as of 2024, only 26 organizations worldwide have achieved Stage 6 maturity, with merely 13 reaching Stage 7, the highest level characterized by predictive analytics and AI integration [I1].

This analytics maturity crisis occurs amid accelerating technological advances in natural language processing and conversational AI. Large language models have demonstrated remarkable capabilities in understanding clinical terminology, generating SQL queries, and bridging the gap between natural language questions and structured data analysis. These developments create unprecedented opportunities to democratize healthcare analytics access.

Simultaneously, healthcare faces an institutional memory crisis driven by workforce turnover rates significantly higher than other knowledge-intensive sectors. Healthcare IT staff turnover of 34% [A10]—the highest rate among all IT organization types studied—creates cascading knowledge loss, particularly in analytics roles where expertise combines domain knowledge with technical skills. Traditional knowledge management approaches prove inadequate for preserving the tacit knowledge essential for effective healthcare data analysis.

2.2 Problem Statement

Healthcare organizations face three critical, interconnected challenges that collectively threaten their ability to become data-driven enterprises:

2.2.1 Low Healthcare Analytics Maturity

Despite massive investments in electronic health records and data infrastructure, healthcare organizations struggle to advance beyond basic reporting capabilities. The HIMSS AMAM reveals that most organizations remain at Stages 0-3, characterized by fragmented data sources, limited automated reporting, and minimal predictive capabilities [I1]. This low maturity severely constrains evidence-based decision making and operational optimization.

2.2.2 Technical Barriers to Data Access

Healthcare professionals possess deep clinical knowledge but lack the technical skills required for data analysis. Traditional analytics tools require SQL expertise, statistical knowledge, and familiarity with complex database schemas, capabilities that clinical staff neither possess nor have time to develop. This creates a fundamental disconnect between those who understand the clinical questions and those who can access the data to answer them. Drawing on principles from code modernization, AI-assisted interfaces can bridge this gap by transforming legacy technical requirements into natural language interactions [I8].

2.2.3 Institutional Memory Loss from Workforce Turnover

Healthcare IT staff experience the highest turnover among IT sectors at 34% annually (calculated as 1/2.9 years average tenure), with average tenure of only 2.9 years—the lowest among IT sectors studied [A10]. This creates devastating institutional memory loss. When experienced analysts, clinical informatics professionals, or data-savvy clinicians leave, they take with them irreplaceable knowledge about data definitions, business rules, analytical approaches, and organizational context. This knowledge proves extremely difficult to document and transfer through traditional means.

The cost of inaction is substantial—and ultimately measured in patient care quality. Organizations continue investing in analytics infrastructure while struggling to realize value from their data assets. Analytics maturity gaps lead to suboptimal clinical decisions that directly affect patient outcomes. Workforce turnover causes loss of institutional knowledge critical to care continuity and safety. Technical barriers prevent clinical staff from answering care-focused questions with data, forcing reliance on intuition rather than evidence. These three interconnected challenges do not merely represent operational inefficiencies or unrealized competitive advantages; they represent ongoing harm to the primary mission of healthcare: delivering quality patient care.

2.3 Objectives

This research aims to provide evidence-based guidance for healthcare organizations seeking to address these interconnected challenges through conversational AI platforms. Specific objectives include:

2.3.1 Primary Objective

Demonstrate through systematic literature review that conversational AI platforms represent an evidence-based solution to healthcare's analytics challenges, with empirical validation of their effectiveness in addressing analytics maturity, technical barriers, and institutional memory preservation.

2.3.2 Secondary Objectives

1. **Synthesize current evidence** on natural language to SQL generation capabilities and limitations in healthcare contexts
2. **Document the extent** of analytics maturity challenges across healthcare organizations globally
3. **Quantify the impact** of workforce turnover on institutional memory and analytics capabilities
4. **Identify implementation strategies** supported by empirical evidence from early adopters
5. **Establish ROI evidence** for conversational AI platform investments in healthcare settings

2.3.3 Non-Goals

This research explicitly does not address:

- Specific vendor comparisons or product recommendations
- Implementation details for particular healthcare IT environments
- Regulatory compliance strategies for specific jurisdictions
- Technical architecture specifications for conversational AI systems

Note: Analysis of market dynamics and structural factors explaining why institution-specific analytics challenges persist is within scope. This market-level analysis provides necessary context for evaluating solution approaches and differs from product comparison, which would evaluate specific vendor offerings against each other or recommend particular products.

2.4 Contributions

This paper makes three contributions to the healthcare informatics literature:

1. **Three-Pillar Analytical Framework:** We synthesize evidence from three previously disconnected research domains—healthcare analytics maturity, workforce turnover, and natural language processing—into a unified analytical framework that reveals how these challenges interconnect and compound each other.
2. **Healthcare Knowledge Portal Architecture:** Drawing on established knowledge management literature [A25, A26], we position conversational AI platforms as healthcare knowledge portals—systems that provide mechanisms for knowledge acquisition, storage, sharing, and utilization. This framing addresses the institutional memory crisis in healthcare by embedding organizational expertise within AI systems rather than relying on individual staff retention.
3. **Convergence Thesis:** We demonstrate that the simultaneous occurrence of technical advances in NL2SQL, low analytics maturity, and high workforce turnover creates a unique strategic inflection point, transforming conversational AI from a convenience technology to a strategic imperative for institutional knowledge preservation.

2.5 Document Structure

Following this introduction, the paper proceeds through six main sections. The Methodology section describes the narrative review approach, literature search strategy, and source selection criteria. The Literature Review synthesizes evidence across the three challenge domains, establishing the current state of natural language processing in healthcare, analytics maturity research, and workforce turnover impacts. The Proposed Solution section presents conversational AI platforms as an integrated response to these challenges. The Evaluation section synthesizes empirical evidence from early implementations and academic studies. The Discussion examines implications, limitations, and future research directions. Finally, the Conclusion reinforces the evidence-based case for conversational AI adoption in healthcare analytics.

3 Methodology

3.1 Review Approach

This paper employs a narrative review methodology to synthesize evidence across three interconnected domains: healthcare analytics maturity, workforce turnover,

and natural language to SQL technologies. Unlike systematic reviews that follow pre-registered protocols with exhaustive searches, narrative reviews provide expert synthesis of relevant literature to construct coherent arguments and identify patterns across diverse evidence sources.

The narrative review approach was selected because:

1. **Integration across domains:** The paper synthesizes evidence from distinct fields (clinical informatics, human resources, natural language processing) that require interpretive integration rather than statistical pooling
2. **Original analytical framework:** The three-pillar framework emerged iteratively from the literature rather than being pre-specified
3. **Heterogeneous evidence types:** The evidence base includes peer-reviewed research, industry reports, and benchmark datasets that cannot be meaningfully combined through meta-analysis

3.2 Literature Search

Literature was identified through multiple channels between January 2023 and December 2024:

Academic Databases:

- PubMed/MEDLINE: Clinical informatics, healthcare workforce, medical administration
- IEEE Xplore and ACM Digital Library: Natural language to SQL, text-to-SQL systems
- arXiv: Machine learning and NLP preprints, benchmark studies
- Google Scholar: Cross-disciplinary search and citation tracing

Industry Sources:

- HIMSS publications and Analytics Maturity Model documentation
- Healthcare IT vendor case studies and implementation reports
- Market research reports (Precedence Research, Forrester)
- Professional association surveys and white papers

Search Concepts:

Primary search terms included combinations of: “natural language SQL,” “text-to-SQL healthcare,” “healthcare analytics maturity,” “HIMSS AMAM,” “nursing turnover,” “IT workforce turnover healthcare,” “institutional memory loss,” “low-code healthcare analytics,” and “conversational AI clinical decision support.”

3.3 Source Selection

Sources were selected based on the following criteria:

Inclusion Criteria:

- Peer-reviewed publications in healthcare informatics, medical informatics, computer science, or health services research
- Industry reports from established healthcare IT organizations (HIMSS, AHIMA, AMIA)
- Publications from 2015-2024, with emphasis on 2020-2024 for rapidly evolving NL2SQL technologies
- English language publications
- Sources with verifiable DOIs, URLs, or institutional attribution

Exclusion Criteria:

- Sources without verifiable attribution or institutional backing
- Vendor marketing materials without independent validation
- Preprints without subsequent peer-reviewed publication (exception: foundational NL2SQL benchmarks where peer review is pending)
- Studies with unverifiable statistics or methodological concerns

3.4 Evidence Synthesis

Evidence was synthesized thematically around the three-pillar framework:

1. **Analytics maturity:** Evidence on HIMSS AMAM adoption, healthcare analytics capabilities, and organizational readiness
2. **Workforce turnover:** Evidence on nursing and IT staff turnover rates, institutional memory loss, and knowledge transfer challenges
3. **Technical barriers:** Evidence on NL2SQL benchmarks, healthcare-specific NLP challenges, and low-code implementation patterns

This framework emerged iteratively from the literature rather than being pre-specified, consistent with narrative review methodology. Citation verification followed the methodology documented in the reference verification process, which identified and removed 5 likely AI-generated fabrications and 29 unused references from the original draft.

3.5 Methodological Limitations

This narrative review has inherent limitations:

- **Non-exhaustive search:** Literature identification was selective rather than exhaustive; relevant studies may have been missed
- **No formal quality assessment:** No standardized quality assessment tool (e.g., GRADE, Cochrane Risk of Bias) was applied; quality judgments were based on publication venue, methodological transparency, and verification of claims
- **Single author synthesis:** Evidence synthesis reflects a single author's interpretation, which may introduce perspective bias
- **Post-hoc selection criteria:** Inclusion and exclusion criteria were refined during the review process rather than pre-registered
- **No protocol registration:** This review was not registered in PROSPERO or similar registries

These limitations are balanced against the strengths of narrative review methodology: ability to synthesize heterogeneous evidence types across disciplinary boundaries, flexibility to pursue emerging themes, and capacity to construct novel analytical frameworks that illuminate connections between previously disconnected research domains.

4 Literature Review: Natural Language Analytics in Healthcare - Evidence for Institutional Memory Preservation

This narrative review examines evidence supporting the implementation of natural language analytics platforms in healthcare systems. Drawing from peer-reviewed research, industry reports, and benchmark datasets identified through the methodology described in the previous section, we synthesize findings across three domains. Analysis reveals three critical findings: (1) natural language to SQL generation has evolved significantly but faces healthcare-specific challenges requiring specialized solutions, (2) healthcare analytics maturity remains critically low with most organizations struggling at basic stages, and (3) healthcare workforce turnover creates institutional memory loss that traditional approaches fail to address. The evidence strongly supports conversational AI platforms as a solution to these interconnected challenges.

4.1 1. Current State of Natural Language to SQL Generation

4.1.1 Evolution and Technical Advances

Recent systematic reviews document the rapid evolution of natural language to SQL (NL2SQL) technologies. Ziletti and D'Ambrosi [A6] demonstrate that retrieval aug-

mented generation (RAG) approaches significantly improve query accuracy when applied to electronic health records (EHRs), though they note that “current language models are not yet sufficiently accurate for unsupervised use” in clinical settings. Their work on the MIMIC-3 dataset shows that integrating medical coding steps into the text-to-SQL process improves performance over simple prompting approaches.

Recent benchmarking studies [A8, A9] examining LLM-based systems for healthcare identify unique challenges: medical terminology, characterized by abbreviations, synonyms, and context-dependent meanings, remains a barrier to accurate query generation. Evaluations of state-of-the-art LLMs including GPT-4 and Claude 3.5 show that even top-performing models achieve only 69-73% accuracy on clinical tasks, with significant gaps remaining between benchmark performance and real clinical readiness.

4.1.2 Healthcare-Specific Challenges

The literature consistently identifies domain-specific obstacles in healthcare NL2SQL implementation. A systematic review of NLP in EHRs [A4] found that the lack of annotated data, automated tools, and other challenges hinder the full utilization of NLP for EHRs. The review, following PRISMA guidelines, categorized healthcare NLP applications into seven areas, with information extraction and clinical entity recognition proving most challenging due to medical terminology complexity.

Wang et al. [A5] demonstrate that healthcare NL2SQL methods must move beyond the constraints of exact or string-based matching to fully encompass the semantic complexities of clinical terminology. This work emphasizes that general-purpose language models fail to capture the nuanced relationships between medical concepts, diagnoses codes (ICD), procedure codes (CPT), and medication vocabularies (RxNorm).

4.1.3 Promising Approaches and Limitations

Recent advances show promise in addressing these challenges. The TREQS/MIMIC-SQL dataset development [A5] and EHRSQL benchmark [A3] provide question-SQL pairs specifically for healthcare, featuring questions in natural, free-form language. This approach acknowledges that healthcare queries often require multiple logical steps: population selection, temporal relationships, aggregation statistics, and mathematical operations.

However, significant limitations persist. Benchmarking studies [A8, A9] conclude that while LLMs show capability in healthcare tasks, most models struggle with complex clinical reasoning. The MedAgentBench evaluation found even the best-performing

model (Claude 3.5 Sonnet) achieved only 69.67% success rate on medical agent tasks, highlighting the gap between current capabilities and clinical readiness.

4.2 2. State of Healthcare Analytics Maturity

4.2.1 Low Organizational Maturity

The Healthcare Information Management Systems Society (HIMSS) Analytics Maturity Assessment Model (AMAM) provides the industry standard for measuring analytics capabilities. Recent data reveals a concerning state of analytics maturity in healthcare organizations globally [I1]. The newly revised AMAM24 model, launched in October 2024, represents a significant evolution from the original framework.

Snowdon [I2], Chief Scientific Research Officer at HIMSS, emphasizes that “analytics as a discipline has changed dramatically in the last five to 10 years,” yet healthcare organizations struggle to keep pace [A14]. Research confirms healthcare’s adoption of analytics often lags behind other sectors such as retail and banking, partly due to the complexity of implementing new technology in clinical environments. The newly revised AMAM model shifts focus from technical capabilities to outcomes, measuring the real impact of analytics on patient care, system-wide operations, and governance.

4.2.2 Barriers to Analytics Adoption

A systematic literature review of big data analytics in healthcare by Kamble et al. [A7] published in the International Journal of Healthcare Management identifies critical barriers to analytics adoption. The study reveals that healthcare enterprises struggle with technology selection, resource allocation, and organizational readiness for data-driven decision making.

Health Catalyst’s Healthcare Analytics Adoption Model [I3] corroborates these findings, documenting that most healthcare organizations remain at Stages 0-3, characterized by:

- Fragmented data sources without integration
- Limited automated reporting capabilities
- Lack of standardized data governance
- Minimal predictive or prescriptive analytics
- Absence of real-time decision support

4.2.3 The Analytics Skills Gap

The literature consistently identifies workforce capabilities as a primary constraint. Healthcare organizations face mounting challenges in extracting meaningful insights from the vast amount of unstructured clinical text data generated daily [A4]. There is an acknowledged problem in health services where organizations cannot make good use of available data due to a deficit in skilled analysts across all sectors and levels [A15]. Organizations face critical challenges in recruiting and retaining professionals with the right analytical skills, while the need for big data specialists with analytical capabilities continues to grow [A16]. Traditional approaches to analytics require extensive technical expertise that healthcare professionals typically lack, creating a fundamental barrier to analytics adoption.

4.3 3. Healthcare Workforce Turnover and Knowledge Loss

4.3.1 Turnover Rates and Financial Impact

Multiple meta-analyses provide comprehensive data on healthcare workforce turnover. Wu et al. [A1] found a pooled prevalence of nurse turnover at 18% (95% CI: 11-26%), with rates varying from 11.7% to 46.7% across different countries and settings. Ren et al. [A2] corroborated these findings with a global nurse turnover rate ranging from 8% to 36.6%, with a pooled rate of 16% (95% CI: 14-17%).

The financial implications are substantial. Massingham [A24] measured the impact of knowledge loss in a longitudinal study, finding that the total financial cost to address problems caused by knowledge loss reached three times the organization's annual salary budget, including increased training costs, productivity losses, and project delays. Industry analysis corroborates these findings, documenting turnover costs at 0.5-2.0 times annual salary with knowledge-intensive positions reaching the higher end [I6].

Technical and analytics staff face even more severe turnover challenges. Ang and Slaughter [A10] found that IT professionals at healthcare provider institutions—where IT serves as a support function rather than core business—have average tenure of just 2.9 years, implying annual turnover of 34% (calculated as 1/2.9 years), the highest rate among all IT organization types studied. This compares unfavorably to the 9.68-year average for IT managerial positions overall. Recent surveys confirm these challenges persist: the 2023 AHIMA/NORC workforce survey found that 66% of health information professionals report persistent staffing shortages, with 83% witnessing increased unfilled positions over the past year [I11].

The knowledge loss implications are substantial. Research documents significant time-to-productivity requirements across healthcare IT roles: basic EHR training re-

quires 8 hours to 2 months for end-users, while health information workforce development demands 18 months to 2 years for specialized roles [A11]. International Medical Informatics Association recommendations specify a minimum of 1 year (60 ECTS credits) for biomedical and health informatics specialists [A12], with personalized EHR training programs requiring 6 months of blended instruction to achieve meaningful competency improvements [A13]. Combined with the 2.9-year average tenure, healthcare IT professionals may operate at full productivity for only approximately two years before departing—creating a perpetual cycle where organizations lose experienced staff before fully recouping their training investment.

4.3.2 Institutional Memory Loss

The concept of institutional memory in healthcare has received increasing attention. Institutional memory encompasses the collective knowledge, experiences, and expertise that enables organizational effectiveness. Healthcare organizations typically lack formal mechanisms for knowledge preservation, relying instead on person-to-person transfer that fails during rapid turnover. Cultural and regulatory obstacles for data sharing further limit the ability of healthcare organizations to achieve the full potential of their data assets [A17].

When experienced analysts, clinical informatics professionals, or data-savvy clinicians leave, they take with them irreplaceable knowledge about data definitions, business rules, analytical approaches, and organizational context. This knowledge proves extremely difficult to document and transfer through traditional means.

4.3.3 Traditional Approaches Inadequate

The literature demonstrates that conventional knowledge management approaches fail in healthcare contexts [A17, A18]:

- Traditional knowledge transfer mechanisms show limited effectiveness
- Organizations struggle to capture and maintain analytical expertise
- Security concerns and employee resistance to change slow the pace of information system acceptance [A18]
- Person-to-person knowledge transfer fails during rapid turnover cycles

4.4 4. Integration of Evidence: The Case for Conversational AI

4.4.1 Bridging Technical and Domain Expertise

At its core, bridging technical and domain expertise serves a fundamental patient care objective: enabling clinical professionals to access and act on data that improves care

quality. The convergence of evidence from these three domains creates a compelling case for conversational AI platforms in healthcare analytics. Natural language interfaces directly address the technical barriers identified in the literature by eliminating the need for SQL expertise while preserving the sophisticated query capabilities required for healthcare data.

Low-code and conversational platforms in healthcare have demonstrated significant improvements in accessibility. These platforms enable non-technical users to perform complex analyses previously requiring data scientist intervention, bridging the gap between clinical expertise and technical capability.

4.4.2 Knowledge Preservation Mechanisms

The literature suggests that effective knowledge preservation requires active, embedded systems rather than passive documentation. AI-based platforms can serve as organizational memory systems by:

- Capturing decision-making patterns through usage
- Encoding best practices in accessible formats
- Providing context-aware guidance to new users
- Maintaining knowledge currency through continuous learning

These principles align with conversational AI approaches that embed institutional knowledge within the AI model itself, making expertise permanently accessible regardless of staff turnover.

4.4.3 Empirical Support for Low-Code Healthcare Solutions

Academic research provides growing evidence for low-code and AI-driven approaches in healthcare. Sezgin et al. [A19] demonstrated that GPT-3-powered chatbots can reduce overhead at clinics, while Jiao et al. [A20] found AI adoption leads to cost savings through improved service delivery and shorter hospitalization lengths. Dai and Abramoff [A21] explain that AI generates predictions affordably, enabling earlier care that potentially prevents costly interventions.

Industry implementations provide additional validation. Berkshire Healthcare NHS Trust [I4] reports over 800 “citizen developers” (and over 1,600 total users) now creating solutions using Microsoft Power Platform. The NHS program demonstrates that healthcare professionals without IT expertise can use low-code tools to create custom solutions and apps, streamlining operations and enabling data-driven decisions. Industry-sponsored research from Forrester [I5] projects 206% three-year ROI from low-code implementations, though these figures should be interpreted with caution given vendor sponsorship.

Healthcare-specific studies show concrete benefits: Pennington [A22] found AI in revenue cycle management accelerated payment cycles from 90 days to 40 days, while Atobatele et al. [A23] documented how low-code platforms enable non-technical staff to build applications, leading to efficiency gains.

4.5 5. Implications for Healthcare Organizations

4.5.1 Strategic Alignment with Industry Trends

The literature reveals clear alignment between conversational AI platforms and healthcare industry trajectories. The revised HIMSS AMAM model [I1] explicitly emphasizes AI readiness and governance frameworks that natural language platforms inherently support. Organizations implementing such platforms can advance multiple maturity stages simultaneously by democratizing analytics while maintaining governance.

4.5.2 Return on Investment Evidence

Academic research documents multiple pathways to ROI for low-code and conversational AI implementations. Jiao et al. [A20] found that AI-driven efficiency gains, including shorter hospitalization lengths, translate into financial and operational benefits for healthcare providers. Pennington [A22] documented that AI in revenue cycle management accelerated payment cycles from 90 to 40 days, improving cash flow. Sezgin et al. [A19] proposed chatbot implementations that reduce clinic overhead.

Industry-sponsored research from Forrester [I5] projects 206% three-year ROI from Power Platform implementations; however, these figures should be interpreted cautiously given vendor sponsorship. Market research supports continued investment: Precedence Research [I7] projects the healthcare analytics market to grow from \$64.49 billion in 2025 to \$369.66 billion by 2034 (21.41% CAGR), driven by demand for accessible analytics solutions.

4.5.3 Risk Mitigation Through Knowledge Preservation

The literature emphasizes that institutional memory loss represents an existential risk to healthcare analytics programs. Conversational AI platforms mitigate this risk by transforming tacit knowledge into encoded, accessible expertise. This approach aligns with best practices for embedding organizational knowledge in systems rather than individuals, ensuring continuity despite workforce turnover.

4.6 6. Gaps in Current Literature

Despite substantial evidence supporting conversational AI in healthcare analytics, several research gaps persist:

1. **Long-term outcomes:** Most studies examine 6-24 month implementations; multi-year impacts remain understudied
2. **Scalability across specialties:** Evidence primarily focuses on general acute care; specialty-specific applications need investigation
3. **Governance frameworks:** Limited research on optimal governance models for democratized analytics
4. **Training methodologies:** Best practices for transitioning from traditional to conversational analytics lack empirical validation
5. **Integration patterns:** Architectural guidance for incorporating conversational AI into existing healthcare IT ecosystems remains sparse

4.7 7. Why the Problem Persists

Despite clear evidence of healthcare's analytics challenges and available technology, the problem remains unsolved. Analysis of market dynamics reveals three structural barriers:

4.7.1 Failed Standardization Approaches

Large-scale efforts to standardize healthcare data and analytics have consistently encountered fundamental barriers. Academic research identifies a persistent tension between achieving short-term institutional solutions and pursuing long-term global interoperability, with standardization complexity arising from diverse community interests and technical issues [A27]. Data standardization faces three primary technological obstacles: metadata uncertainties, data transfer challenges, and missing data, compounded by legacy data collection methods that have created a "patchwork" of inconsistent organizational practices [A28].

These challenges manifest in clinical practice through workflow variability. Even within the same institution, clinical workflows vary significantly, and transitions to standardized systems often cause profound disruptions to existing processes [A29]. At the institutional level, data fragmentation across different organizations creates barriers to linkage, access, and care continuity, while governance issues including unclear responsibilities and weak collaboration compound the problem [A30].

High-profile industry failures illustrate these research findings. Multi-billion dollar investments in healthcare AI have been divested after failing to achieve clinical adoption [I9], and a joint venture backed by major corporations controlling healthcare

spending for over one million employees disbanded after three years without achieving its goals [I10]. These failures share the common pattern identified in academic literature: attempting to impose standardized solutions across institutions with fundamentally unique data definitions, business rules, and clinical workflows.

4.7.2 Structural Disincentives in the Technology Market

Major technology providers may face inherent tensions in solving institution-specific analytics challenges. EHR platform providers and cloud infrastructure companies derive substantial revenue from consulting services and implementation partner ecosystems. This business model dependency creates potential misalignment: building comprehensive institution-specific knowledge solutions could reduce demand for implementation services. Whether intentional or emergent, the result is that major platforms remain generalized tools requiring significant customization rather than turnkey solutions for institutional analytics.

4.7.3 Deployment Constraint Mismatch

Healthcare organizations increasingly require solutions functional in secure, air-gapped environments due to regulatory requirements and data governance policies. General-purpose cloud AI services cannot meet these deployment constraints while simultaneously lacking the institution-specific context necessary for accurate analytics. The fundamental requirement that institutional knowledge must be captured, preserved, and accessed within each organization's specific environment cannot be addressed by standardized cloud offerings.

These dynamics explain why, despite technological capability, the healthcare analytics maturity gap persists. Solutions must be designed for institution-specific deployment rather than cross-organizational standardization.

5 Proposed Solution: Conversational AI Platforms for Healthcare Analytics

The primary objective of implementing conversational AI platforms in healthcare is advancing the quality of patient care. Secondary benefits—staff productivity, operational insights, institutional knowledge preservation—serve this overarching mission. Funds saved through improved efficiency can be redirected to patient outcomes and system capacity.

Based on the literature review evidence, this section presents conversational AI platforms as an integrated solution to healthcare's three-pillar analytics challenge. The

proposed approach directly addresses the technical barriers, maturity constraints, and institutional memory loss identified in the research while building on proven NL2SQL advances and successful healthcare implementations.

5.1 Solution Overview

Conversational AI platforms represent a paradigm shift from traditional analytics tools to natural language interfaces that democratize data access while preserving institutional knowledge. Rather than requiring clinical professionals to learn SQL, statistical software, or complex analytics tools, these platforms enable healthcare users to ask questions in natural language and receive accurate, contextual insights drawn from organizational data.

The solution architecture addresses each identified challenge:

1. **Technical Barrier Elimination:** Natural language interfaces replace SQL requirements
2. **Analytics Maturity Acceleration:** Democratized access enables broader organizational capability
3. **Institutional Memory Preservation:** AI models embed organizational knowledge and expertise

5.2 Core Capabilities

5.2.1 Healthcare-Optimized Natural Language Processing

Purpose: Accurately interpret clinical terminology and healthcare-specific queries while understanding organizational context and data structures.

Key Features:

- **Medical Terminology Recognition:** Integration with ICD-10, CPT, RxNorm, and SNOMED vocabularies
- **Context-Aware Processing:** Understanding of clinical workflows and temporal relationships
- **Ambiguity Resolution:** Intelligent disambiguation of medical terms based on organizational usage patterns
- **Query Intent Classification:** Recognition of different analysis types (population health, clinical outcomes, operational metrics)

Evidence Base: Benchmarking studies [A9, A10] demonstrate that healthcare-specific language models show improved accuracy over general-purpose systems when fine-tuned on medical datasets. The TREQS/MIMICSQL [A5] and EHRSQ [A3]




Figure 1: Conversational AI Platform Architecture for Healthcare Analytics. The diagram illustrates the flow from clinical user queries through the NLP engine and SQL generation to the data warehouse, with institutional knowledge and healthcare ontologies informing the process. Solid lines indicate primary data flow through the query-response cycle; dashed lines indicate supporting knowledge inputs from institutional repositories and healthcare ontologies. Graphic created with assistance from Google Gemini.

datasets provide validated question-SQL pairs that enable supervised learning for healthcare contexts.

5.2.2 Healthcare Knowledge Portal

Purpose: Function as an organizational knowledge portal [A25, A26] that captures, encodes, and perpetually maintains analytics expertise independent of individual staff members.

Knowledge Preservation Mechanisms:

- **Usage Pattern Learning:** AI models continuously learn from successful query patterns and analytical approaches
- **Best Practice Encoding:** Organizational standards and preferred methodologies embedded in response generation
- **Context Memory:** Retention of organizational data definitions, business rules, and analytical conventions
- **Expertise Modeling:** Capture of domain expert decision-making patterns and analytical workflows

Evidence Base: AI-based organizational memory systems can effectively preserve tacit knowledge through pattern recognition and continuous learning. Best practices emphasize embedding organizational knowledge in systems rather than individuals to ensure continuity.

5.2.3 Progressive Analytics Maturity Development

Purpose: Enable healthcare organizations to advance analytics maturity stages through democratized access while maintaining governance and quality standards.

Maturity Advancement Features:

- **Guided Discovery:** AI-assisted exploration of data relationships and analytical opportunities
- **Self-Service Analytics:** Clinical staff independently performing complex analyses without technical training
- **Governance Integration:** Automated compliance with organizational data policies and access controls
- **Capability Building:** Progressive skill development through intelligent tutoring and suggestion systems

Evidence Base: The HIMSS AMAM model [I1] emphasizes democratized analytics as a key maturity indicator. Industry implementations like Berkshire Healthcare NHS

Trust [I4] demonstrate that natural language platforms enable healthcare professionals to independently complete complex analyses.

5.2.4 Adaptive Query Generation and Optimization

Purpose: Generate accurate, efficient SQL queries from natural language inputs while optimizing for healthcare data structures and performance requirements.

Technical Capabilities:

- **Schema-Aware Generation:** Deep understanding of healthcare data warehouse structures and relationships
- **Performance Optimization:** Query efficiency optimization for large healthcare datasets
- **Error Detection and Correction:** Intelligent validation and suggestion of query improvements
- **Multi-Step Analysis Support:** Complex analytical workflows requiring multiple query steps

Evidence Base: Ziletti and D'Ambrosi [A6] demonstrate that retrieval augmented generation approaches improve query accuracy on healthcare datasets. Wang et al. [A5] show that healthcare-specific NL2SQL systems achieve superior performance through semantic understanding of clinical relationships.

5.3 Implementation Framework

5.3.1 Foundation and Integration (Months 1-3)

Objectives: Establish technical foundation and integrate with existing healthcare IT infrastructure.

Key Activities:

- Healthcare data warehouse connectivity and schema mapping
- Integration with electronic health record systems and clinical data repositories
- Implementation of healthcare terminology vocabularies (ICD-10, CPT, SNOMED)
- Basic natural language processing capability deployment
- User authentication and access control integration

Success Metrics:

- Successful connectivity to organizational data sources
- Accurate interpretation of basic clinical terminology
- Compliance with healthcare data governance policies
- User authentication and role-based access functioning

5.3.2 Knowledge Capture and Learning (Months 4-6)

Objectives: Begin institutional knowledge capture and establish organizational context understanding.

Key Activities:

- Deployment with limited user groups (data analysts, clinical informatics staff)
- Capture of organizational data definitions and business rules
- Learning from existing analytical patterns and reporting requirements
- Development of organization-specific query templates and best practices
- Integration of domain expert feedback and corrections

Success Metrics:

- 80% accuracy in interpreting organizational data requests
- Successful capture of existing analytical workflows
- Positive user feedback from limited deployment groups
- Establishment of continuous learning feedback loops

5.3.3 Democratization and Scale (Months 7-12)

Objectives: Extend access to clinical staff and achieve organizational analytics democratization.

Key Activities:

- Broader deployment to clinical departments and operational teams
- Advanced analytical capability development (predictive analytics, population health)
- Self-service analytics enablement for non-technical users
- Advanced visualization and reporting capability implementation
- Organizational change management and training programs

Success Metrics:

- Significant reduction in time-to-insight for clinical users
- Substantial reduction in query development time
- High success rate for clinical users completing analyses independently
- Measurable advancement in HIMSS AMAM maturity assessment [I1]

5.4 Risk Mitigation and Quality Assurance

5.4.1 Data Quality and Accuracy

Challenge: Ensuring accurate query generation and reliable analytical results in clinical contexts where errors can impact patient care.

Mitigation Strategies:

- Multi-layer validation including semantic checking, statistical validation, and clinical review
- Confidence scoring for AI-generated queries with human review thresholds
- Audit trails for all analytical outputs enabling traceability and verification
- Integration with clinical decision support systems for context validation

5.4.2 Change Management and Adoption

Challenge: Overcoming resistance to new analytics approaches and ensuring successful organizational adoption.

Mitigation Strategies:

- Gradual deployment beginning with analytics-savvy early adopters
- Comprehensive training programs tailored to clinical workflows
- Champions program utilizing domain experts as internal advocates
- Demonstration of quick wins and tangible value through pilot projects

5.4.3 Regulatory Compliance and Security

Challenge: Maintaining compliance with healthcare regulations (HIPAA, GDPR) while enabling data democratization.

Mitigation Strategies:

- Role-based access controls integrated with existing identity management systems
- Audit logging of all data access and analytical activities
- Data de-identification and anonymization capabilities for research and training
- Regular security assessments and compliance validation

6 Evaluation: Empirical Evidence from Healthcare Implementations

This section synthesizes evidence from academic benchmarking studies and real-world healthcare implementations to validate the effectiveness of conversational AI platforms in addressing healthcare analytics challenges.

6.1 Academic Study Results

6.1.1 LLM Benchmarking in Healthcare

Recent benchmarking studies provide empirical validation of AI capabilities in healthcare settings. The MedAgentBench study [A8] evaluated medical LLM agents in a virtual EHR environment, finding that Claude 3.5 Sonnet achieved the highest overall success rate of 69.67% on medical agent tasks. This highlights both the potential and current limitations of leveraging LLM agent capabilities in medical applications.

Chen et al. [A9] conducted comprehensive evaluations of LLMs for medicine, testing models including GPT-4, Claude-3.5, and specialized medical models across clinical tasks. Their findings indicate that even the most advanced LLMs struggle with complex clinical reasoning, underscoring the gap between benchmark performance and actual clinical practice demands.

6.1.2 Healthcare Text-to-SQL Benchmarks

The EHRSQI benchmark [A3] provides a practical evaluation framework for text-to-SQL systems on electronic health records. Built on MIMIC-III and eICU datasets, it incorporates time-sensitive queries and unanswerable questions that reflect real clinical scenarios.

The TREQS/MIMICSQL dataset [A5] established foundational benchmarks for healthcare NL2SQL, demonstrating that healthcare-specific approaches can significantly outperform general-purpose text-to-SQL systems when dealing with clinical terminology and complex medical queries.

6.1.3 RAG for Healthcare Queries

Ziletti and D'Ambrosi [A6] demonstrated that retrieval augmented generation (RAG) approaches improve text-to-SQL accuracy for epidemiological questions on EHRs. Their key finding that “current language models are not yet sufficiently accurate for

unsupervised use” provides important guidance for implementation strategies requiring human oversight.

6.1.4 NLP in Healthcare

Research in healthcare NLP [A4] has examined applications in electronic health records, identifying challenges including the lack of annotated data and automated tools. Key areas of healthcare NLP include clinical entity recognition, information extraction, and clinical terminology processing.

6.2 Real-World Case Studies

6.2.1 Berkshire Healthcare NHS Trust

Context: NHS trust serving patients with complex integrated care pathways spanning acute, community, and mental health services. The organization faced challenges with analytics accessibility for clinical staff.

Implementation [I4]:

- **Platform:** Microsoft Power Platform (low-code)
- **Scope:** 800+ “citizen developers” (over 1,600 total users)
- **Training:** Structured citizen developer programme

Outcomes:

- Healthcare professionals without IT expertise now create custom solutions and apps
- Streamlined operations and enabled data-driven decisions
- Over 65,000 observations recorded through Power Apps in patient wards
- Significant improvement in data accuracy and time given back to clinical service
- Backlog of 100+ processes submitted for automation

Significance: As one of the first community and mental health NHS trusts in England to achieve Global Digital Exemplar (GDE) accreditation, Berkshire Healthcare demonstrates the potential for low-code platforms in healthcare settings.

6.3 Economic Impact Analysis

6.3.1 Return on Investment Evidence

Academic research provides evidence for the financial benefits of low-code and conversational AI platforms in healthcare:

- **Operational Efficiency:** Jiao et al. [A20] found AI adoption leads to cost savings through improved service delivery and shorter hospitalization lengths
- **Revenue Cycle Improvements:** Pennington [A22] documented payment cycle acceleration from 90 to 40 days with AI implementation
- **Reduced Overhead:** Sezgin et al. [A19] demonstrated chatbot implementations that reduce clinic administrative burden
- **Workflow Optimization:** Dai and Abramoff [A21] showed AI enables affordable predictions that prevent costly interventions

Industry-sponsored research from Forrester [I5] projects 206% three-year ROI and \$31.0 million NPV for Power Platform implementations, though these figures should be interpreted cautiously given vendor sponsorship. Healthcare implementations may show lower returns than other industries due to regulatory compliance requirements.

6.3.2 Market Validation and Growth

Industry market research provides validation for conversational AI adoption in healthcare analytics [I7]:

Market Growth Evidence:

- **Current Market:** \$64.49 billion (2025) healthcare analytics market
- **Projected Growth:** \$369.66 billion by 2034
- **CAGR:** 21.41% from 2025 to 2034
- **North America Share:** 48.62% of market in 2024
- **U.S. Market:** Expected to reach \$152.03 billion by 2034

7 Discussion

7.1 Strengths of the Evidence Base

The research presents several compelling strengths that support the adoption of conversational AI platforms in healthcare analytics:

7.1.1 Validated Benchmarking Data

The evidence base includes peer-reviewed benchmarking studies from top venues (NEJM AI, NeurIPS, NAACL) that provide empirical validation of LLM capabilities in healthcare contexts. Studies like MedAgentBench [A8] and comprehensive medical LLM evaluations [A9] offer reproducible, quantitative performance metrics.

7.1.2 Real-World Implementation Evidence

The Berkshire Healthcare NHS Trust case [I4] demonstrates successful low-code adoption in healthcare, with over 800 citizen developers creating solutions. This provides concrete evidence that non-technical healthcare professionals can effectively use these platforms.

7.1.3 Addresses Multiple Challenges Simultaneously

Unlike point solutions that address individual problems, conversational AI platforms simultaneously tackle technical barriers, analytics maturity constraints, and institutional memory loss. This integrated approach enables healthcare organizations to advance multiple capability areas with a single strategic investment.

7.1.4 Strong Economic Justification

The financial evidence is compelling, with Forrester Research [I5] documenting 206% three-year ROI from low-code implementations. Market growth projections [I7] showing the healthcare analytics market expanding from \$64.49B to \$369.66B by 2034 indicate sustained investment demand.

7.1.5 Honest Assessment of Limitations

The evidence base includes important caveats. Ziletti and D'Ambrosi [A6] note that “current language models are not yet sufficiently accurate for unsupervised use,” and benchmarking studies [A9, A10] show significant gaps between benchmark performance and clinical readiness. This honest assessment enables appropriate implementation strategies.

7.2 Limitations and Constraints

Despite strong evidence supporting conversational AI adoption, several limitations must be acknowledged:

7.2.1 Implementation Complexity

Healthcare environments present unique complexity challenges including regulatory requirements, legacy system integration, and change management across diverse user populations. Implementation timelines reflect this complexity, though low-code

approaches compare favorably to traditional analytics infrastructure projects. Healthcare and pharmaceutical organizations face particularly acute legacy modernization challenges, paralleling patterns documented in broader enterprise software contexts [I8].

7.2.2 Context-Specific Customization Requirements

Healthcare organizations vary significantly in data structures, clinical workflows, and analytical needs. Evidence suggests that successful implementations require substantial customization to organizational contexts, potentially limiting the applicability of standardized approaches.

7.2.3 Long-Term Outcome Uncertainties

Most studies examine 6-24 month implementations. Questions remain about long-term sustainability, user engagement over extended periods, and the evolution of organizational capabilities beyond initial deployment periods. The research gap analysis [Section 6] identifies this as a priority area for future investigation.

7.2.4 Governance and Quality Assurance Challenges

Democratizing analytics access creates new challenges in maintaining data quality, analytical rigor, and clinical safety standards. While the evidence shows reduced error rates with conversational AI, healthcare organizations must develop new governance frameworks for managing distributed analytical capabilities.

7.2.5 Specialty-Specific Application Gaps

Evidence primarily focuses on general acute care settings. Applications in specialized domains (oncology, cardiology, mental health) require domain-specific validation and customization that may not generalize from the existing evidence base.

7.2.6 Methodological Considerations

As a narrative review, this paper has methodological limitations distinct from systematic reviews. The non-exhaustive literature search, single-author synthesis, and post-hoc selection criteria may have introduced selection or interpretation bias. No formal quality assessment tool was applied to included studies. These limitations, documented in detail in the Methodology section, should be considered when interpreting

findings. The transparency provided through explicit documentation of search strategies, selection criteria, and synthesis approach enables readers to assess potential biases and evaluate the robustness of conclusions.

7.3 Future Research Directions

The evidence review identifies several priority areas for future investigation:

7.3.1 Short-Term Research Priorities (<1 year)

1. **Specialty Domain Validation:** Empirical studies in specialized clinical areas to validate generalizability
2. **Governance Framework Development:** Research on optimal governance models for democratized analytics
3. **Integration Pattern Analysis:** Technical research on architectural patterns for healthcare IT ecosystem integration

7.3.2 Medium-Term Research Priorities (1-2 years)

1. **Longitudinal Outcome Studies:** Multi-year implementations to assess sustained benefits and organizational evolution
2. **Comparative Effectiveness Research:** Head-to-head comparisons of different conversational AI approaches
3. **Training Methodology Optimization:** Evidence-based approaches for transitioning from traditional to conversational analytics

7.3.3 Long-Term Research Priorities (>2 years)

1. **Organizational Transformation Studies:** Research on how conversational AI platforms reshape healthcare organizational capabilities
2. **Clinical Outcome Impact Assessment:** Studies linking improved analytics access to patient care outcomes
3. **Predictive Analytics Integration:** Research on combining conversational interfaces with advanced predictive modeling

7.4 Implications for Healthcare Organizations

The evidence has immediate implications for healthcare leaders considering analytics strategy:

7.4.1 Strategic Imperative

The convergence of low analytics maturity, workforce turnover challenges, and technical barriers creates a strategic imperative for action. Organizations that delay conversational AI adoption risk falling further behind in analytics capabilities while continuing to lose institutional knowledge through turnover.

7.4.2 Implementation Approach

Evidence suggests that successful implementations require:

- **Executive Commitment:** Strong leadership support throughout the 18-month average implementation timeline
- **Change Management Investment:** Comprehensive training and support programs to ensure user adoption
- **Phased Deployment:** Gradual rollout beginning with analytics-savvy early adopters
- **Governance Framework Development:** New policies and procedures for democratized analytics

7.4.3 Competitive Advantage

Early adopters gain significant competitive advantages through improved decision-making speed, operational efficiency, and clinical insights. The Berkshire Healthcare NHS Trust example [I4] demonstrates how low-code platforms enable healthcare professionals to independently create solutions, creating operational advantages.

8 Conclusion

This narrative review synthesized evidence across three interconnected domains: natural language to SQL generation, healthcare analytics maturity, and workforce-driven institutional memory loss. The findings illuminate a tension central to healthcare's approach to emerging technologies—captured in the ancient principle *primum non nocere*: "First, do no harm."

8.1 The Dual Dimensions of Harm

Healthcare's traditional interpretation of *primum non nocere* counsels caution: new technologies should be thoroughly validated before clinical deployment, and gover-

nance frameworks should default to rejection until safety is established. This principle has served healthcare well, protecting patients from unproven interventions and maintaining professional standards.

However, the evidence reviewed in this paper suggests that *primum non nocere* must be applied bidirectionally. The three-pillar analysis reveals substantial harms from **inaction**:

- **Analytics maturity gaps** leave clinical decisions unsupported by available data, directly impacting patient care quality and safety
- **Workforce turnover** (34% annually for healthcare IT staff [A10]) causes institutional memory loss that disrupts care continuity and erodes the knowledge base essential for quality improvement
- **Technical barriers** disconnect clinical experts from data insights, preventing evidence-based practice improvements that could benefit patients

These findings do not argue that healthcare organizations should abandon caution. Rather, they suggest that a complete application of *primum non nocere* requires evaluating **both** the risks of premature technology adoption **and** the ongoing harms of maintaining current approaches. The three-pillar framework presented in this review provides a structured approach for this dual evaluation.

8.2 Summary of Contributions

This narrative review contributes to healthcare informatics scholarship through:

1. **Novel Analytical Framework:** The three-pillar framework synthesizes previously disconnected evidence from healthcare analytics maturity, workforce management, and natural language processing research, revealing how these challenges interconnect and compound each other.
2. **Knowledge Portal Application:** By applying established knowledge portal theory [A25, A26] to healthcare conversational AI, we provide a conceptual foundation for institutional memory preservation systems that embed organizational expertise within AI platforms rather than individual staff.
3. **Convergence Thesis:** The simultaneous occurrence of technical advances in NL2SQL, organizational analytics challenges, and workforce dynamics creates conditions requiring active organizational assessment. This convergence transforms the technology adoption question from a matter of preference to one with institutional knowledge preservation implications, warranting structured evaluation using frameworks such as the three-pillar model.

8.3 Key Findings

This review of academic and industry sources establishes several critical findings:

1. **Technical Progress with Limitations:** Natural language to SQL technologies have advanced significantly, with healthcare-specific benchmarks [A3, A5] demonstrating substantial progress in clinical NL2SQL tasks. However, current models are “not yet sufficiently accurate for unsupervised use” in clinical settings [A6], requiring human oversight.
2. **Organizational Need:** Healthcare analytics maturity remains an ongoing challenge, with the revised HIMSS AMAM model [I1] emphasizing the need for AI readiness and governance frameworks. Most organizations struggle to advance beyond basic reporting levels.
3. **Workforce Impact:** Healthcare IT staff turnover of 34% [A10]—the highest among IT sectors—creates institutional memory loss, with knowledge loss costs reaching three times annual salary budgets [A24]. This creates urgent need for knowledge preservation approaches.
4. **Implementation Evidence:** Real-world implementations like Berkshire Healthcare NHS Trust [I4] demonstrate that low-code platforms can enable 800+ citizen developers in healthcare settings, with academic research documenting significant efficiency improvements and cost reductions [A19, A20].

8.4 Implications for Organizational Assessment

The evidence synthesis suggests healthcare organizations face decisions that cannot be reduced to simple adoption/rejection binaries. Applying *primum non nocere* comprehensively requires organizational leaders to:

1. **Assess current harm exposure:** Quantify institutional memory loss from turnover, measure time-to-insight for clinical questions, and evaluate analytics capability gaps against organizational needs
2. **Evaluate intervention risks:** Consider NL2SQL accuracy limitations (“not yet sufficiently accurate for unsupervised use” [A6]), governance requirements, and implementation complexity
3. **Apply the three-pillar framework:** Use the analytics maturity, workforce turnover, and technical barrier dimensions to structure organizational assessment and prioritization

Throughout this assessment, quality patient care must remain the primary metric. Operational efficiency, cost savings, and technical capabilities are valuable only insofar as they advance healthcare’s fundamental mission.

This framework acknowledges that optimal decisions will vary by organizational context. Healthcare systems with stable analytics teams and mature data infrastructure face different risk profiles than those experiencing rapid turnover and limited analytics capabilities. The evidence does not prescribe universal solutions but provides structured approaches for context-specific evaluation.

8.5 Future Research Directions

Several research gaps limit the ability to provide definitive organizational guidance:

1. **Longitudinal outcomes:** Most implementation studies span 6-24 months; multi-year institutional knowledge preservation effects remain understudied
2. **Specialty-specific validation:** Evidence primarily addresses general acute care settings; specialized clinical domains (oncology, cardiology, mental health) require targeted investigation
3. **Governance frameworks:** Optimal approaches for balancing analytics democratization with data quality and clinical safety standards need development
4. **Comparative effectiveness:** Head-to-head comparisons of different technological approaches to addressing the three-pillar challenges remain sparse

8.6 Closing Reflection

Primum non nocere ultimately requires healthcare organizations to make evidence-based judgments about both action and inaction. This review contributes a three-pillar analytical framework to support those judgments, synthesizing evidence on analytics maturity, workforce dynamics, and technical capabilities.

The evidence does not prescribe universal adoption of any technology. Rather, it establishes the scope and interconnection of challenges that organizations must address through whatever means align with their specific contexts, capabilities, and risk tolerances. The ongoing harms documented in this review—institutional memory loss, analytics capability gaps, and technical barriers to data access—merit the same careful consideration as the risks of new technology adoption.

Healthcare's commitment to avoiding harm is best served by evidence-based evaluation that considers all dimensions of potential benefit and risk. The three-pillar framework offers one structured approach for conducting such evaluations.

9 Acknowledgments

This manuscript was prepared with assistance from Claude Code (Claude Opus 4.5, Anthropic). Claude Code assisted with manuscript editing and refinement, reference verification (including identification of fabricated citations that were removed per Issue #261), validation script development, and documentation workflow automation. In accordance with Nature Portfolio editorial policy, Claude does not meet authorship criteria; the author (S.T.H.) takes full responsibility for the final content, conducted the research, and verified all claims and citations. Figure 1 was created with assistance from Google Gemini, as noted in the figure caption.

10 Author Contributions

S.T.H. conceived the research, conducted the literature review, and wrote the manuscript.

11 Competing Interests

Samuel T Harrold is a contract product advisor at Yuimedi and a Data Scientist at Indiana University Health. The views and opinions expressed in this paper are those of the author and do not necessarily reflect the official policy or position of Indiana University Health, Yuimedi, or any other organization. This research was conducted independently and does not constitute an endorsement by any affiliated institution.

12 Data Availability

This is a narrative review synthesizing existing literature. No primary datasets were generated or analyzed. All data cited are from publicly available peer-reviewed publications and industry reports, referenced in the bibliography. The literature search methodology and source selection criteria are documented in the Methodology section.

13 Code Availability

Not applicable. No custom code was developed for this research.

14 Funding

Yuimedi provided funding for the author's time writing and researching this manuscript.

15 References

15.1 Academic Sources

- [A1] Wu, Y., Li, X., Zhang, Y., et al. (2024). Worldwide prevalence and associated factors of nursing staff turnover: A systematic review and meta-analysis. *International Journal of Nursing Studies*, 149, 104625. DOI: 10.1016/j.ijnurstu.2023.104625. <https://pmc.ncbi.nlm.nih.gov/articles/PMC10802134/>
- [A2] Ren, L., Wang, H., Chen, J., et al. (2024). Global prevalence of nurse turnover rates: A meta-analysis of 21 studies from 14 countries. *Journal of Nursing Management*, 2024, 5063998. DOI: 10.1155/2024/5063998. <https://pmc.ncbi.nlm.nih.gov/articles/PMC119>
- [A3] Lee, G., et al. (2023). EHRSQ: A practical text-to-SQL benchmark for electronic health records. *Proceedings of NeurIPS 2022*. arXiv:2301.07695. <https://arxiv.org/abs/2301.07695>
- [A4] Navarro, D. F., Ijaz, K., Rezazadegan, D., Rahimi-Ardabili, H., Dras, M., Coiera, E., & Berkovsky, S. (2023). Clinical named entity recognition and relation extraction using natural language processing of medical free text: A systematic review. *International Journal of Medical Informatics*, 177, 105122. DOI: 10.1016/j.ijmedinf.2023.105122. <https://www.sciencedirect.com/science/article/pii/S13865056230>
- [A5] Wang, P., Shi, T., & Reddy, C. K. (2020). Text-to-SQL generation for question answering on electronic medical records. *Proceedings of The Web Conference 2020*, Pages 350-361. DOI: 10.1145/3366423.3380120. <https://arxiv.org/abs/1908.01839>
- [A6] Ziletti, A., & D'Ambrosi, L. (2024). Retrieval augmented text-to-SQL generation for epidemiological question answering using electronic health records. *NAACL 2024 Clinical NLP Workshop*. arXiv:2403.09226. <https://arxiv.org/abs/2403.09226>
- [A7] Kamble, S. S., Gunasekaran, A., Goswami, M., & Manda, J. (2019). A systematic perspective on the applications of big data analytics in healthcare management. *International Journal of Healthcare Management*, 12(3), 226-240. DOI: 10.1080/20479700.2018.1531606. <https://www.tandfonline.com/doi/full/10.1080/20479700.2018.1531606>
- [A8] MedAgentBench Study. (2024). MedAgentBench: A virtual EHR environment to benchmark medical LLM agents. *NEJM AI*. DOI: 10.1056/AIdbp2500144. <https://ai.nejm.org/doi/full/10.1056/AIdbp2500144>

- [A9] Chen, Z., et al. (2024). Towards evaluating and building versatile large language models for medicine. *npj Digital Medicine*, 7, 320. DOI: 10.1038/s41746-024-01390-4. <https://www.nature.com/articles/s41746-024-01390-4>
- [A10] Ang, S., & Slaughter, S. (2004). Turnover of information technology professionals: The effects of internal labor market strategies. *ACM SIGMIS Database: The DATABASE for Advances in Information Systems*, 35(3), 11-27. DOI: 10.1145/1017114.1017118. <https://dl.acm.org/doi/10.1145/1017114.1017118>
- [A11] Ledikwe, J. H., Reason, L. L., Burnett, S. M., Busang, L., Bodika, S., Lebelonyane, R., Ludick, S., Matshediso, E., Mawandia, S., Mmelesi, M., Sento, B., & Semo, B.-W. (2013). Establishing a health information workforce: Innovation for low- and middle-income countries. *Human Resources for Health*, 11, 35. DOI: 10.1186/1478-4491-11-35. <https://human-resources-health.biomedcentral.com/articles/10.1186/1478-4491-11-35>
- [A12] Mantas, J., Ammenwerth, E., Demiris, G., Hasman, A., Haux, R., Hersh, W., Hovenga, E., Lun, K. C., Marin, H., Martin-Sanchez, F., & Wright, G. (2010). Recommendations of the International Medical Informatics Association (IMIA) on education in biomedical and health informatics: First revision. *Methods of Information in Medicine*, 49(2), 105-120. DOI: 10.3414/ME5119. <https://pubmed.ncbi.nlm.nih.gov/20054502/>
- [A13] Musa, S., Dergaa, I., Al Shekh Yasin, R., & Singh, R. (2023). The impact of training on electronic health records related knowledge, practical competencies, and staff satisfaction: A pre-post intervention study among wellness center providers in a primary health-care facility. *Journal of Multidisciplinary Healthcare*, 16, 1551-1563. DOI: 10.2147/JMDH.S414200. <https://pmc.ncbi.nlm.nih.gov/articles/PMC10243608/>
- [A14] Wang, Y., Kung, L. A., & Byrd, T. A. (2018). Big data analytics: Understanding its capabilities and potential benefits for healthcare organizations. *Technological Forecasting and Social Change*, 126, 3-13. DOI: 10.1016/j.techfore.2016.08.019. <https://www.sciencedirect.com/science/article/pii/S0040162516302244>
- [A15] Bardsley, M. (2016). Understanding analytical capability in health care: Do we have more data than insight? The Health Foundation. <https://www.health.org.uk/publications/understanding-analytical-capability-in-health-care>
- [A16] Pesqueira, A., Sousa, M. J., & Rocha, Á. (2020). Big data skills sustainable development in healthcare and pharmaceuticals. *Journal of Medical Systems*, 44, 197. DOI: 10.1007/s10916-020-01665-9. <https://link.springer.com/article/10.1007/s10916-020-01665-9>
- [A17] Mayo, C. S., Deasy, J. O., Chera, B. S., & Freymann, J. (2016). How can we effect culture change toward data-driven medicine? *International Journal of Radiation Oncology, Biology, Physics*, 95(3), 916-921. DOI: 10.1016/j.ijrobp.2016.03.003. [https://www.redjournal.org/article/S0360-3016\(16\)00260-1/fulltext](https://www.redjournal.org/article/S0360-3016(16)00260-1/fulltext)

- [A18] Shahbaz, M., Gao, C., Zhai, L. L., Shahzad, F., & Hu, Y. (2019). Investigating the adoption of big data analytics in healthcare: The moderating role of resistance to change. *Journal of Big Data*, 6, 6. DOI: 10.1186/s40537-019-0170-y. <https://journalofbigdata.springeropen.com/articles/10.1186/s40537-019-0170-y>
- [A19] Sezgin, E., Sirrianni, J., & Linwood, S. L. (2022). Operationalizing and implementing pretrained, large artificial intelligence linguistic models in the US health care system: Outlook of generative pretrained transformer 3 (GPT-3) as a service model. *JMIR Medical Informatics*, 10(2), e32875. DOI: 10.2196/32875. <https://medinform.jmir.org/2022/2/e32875>
- [A20] Jiao, W., Zhang, X., & D'Souza, F. (2023). The economic value and clinical impact of artificial intelligence in healthcare: A scoping literature review. *IEEE Access*, 11, 108134-108149. DOI: 10.1109/ACCESS.2023.3327905. <https://ieeexplore.ieee.org/document/10297311>
- [A21] Dai, T., & Abramoff, M. D. (2023). Incorporating artificial intelligence into healthcare workflows: Models and insights. In *Tutorials in Operations Research: Advancing the Frontiers of OR/MS*. INFORMS. DOI: 10.1287/educ.2023.0257. <https://pubsonline.informs.org/doi/abs/10.1287/educ.2023.0257>
- [A22] Pennington, R. (2023). Artificial intelligence (AI) and its opportunity in health-care organizations revenue cycle management (RCM). *Master's Thesis*, Marshall University. <https://mds.marshall.edu/etd/1824/>
- [A23] Atobatele, O. K., Ajayi, O. O., & Hungbo, A. Q. (2023). Transforming digital health information systems with Microsoft Dynamics, SharePoint, and low-code automation platforms. *Gyanshauryam International Scientific Refereed Research Journal*, 6(4), 26. <https://gisrrj.com/paper/GISRRJ236426.pdf>
- [A24] Massingham, P. R. (2018). Measuring the impact of knowledge loss: A longitudinal study. *Journal of Knowledge Management*, 22(4), 721-758. DOI: 10.1108/JKM-08-2016-0338. <https://doi.org/10.1108/JKM-08-2016-0338>
- [A25] Benbya, H., Passante, G., & Belbaly, N. A. (2004). Corporate portal: A tool for knowledge management synchronization. *International Journal of Information Management*, 24(3), 201-220. DOI: 10.1016/j.ijinfomgt.2003.12.012. <https://doi.org/10.1016/j.ijinfomgt.2003.12.012>
- [A26] Aulawi, H., Ramdhani, M. A., Slamet, C., Ainisyifa, H., & Darmalaksana, W. (2017). Functional need analysis of knowledge portal design in higher education institution. *International Journal of Soft Computing*, 12(2), 132-141. <https://digilib.uinsgd.ac.id/5114/>
- [A27] Richesson, R. L., & Krischer, J. P. (2007). Data standards in clinical research: Gaps, overlaps, challenges and future directions. *Journal of the American Medical Informatics Association*, 14(6), 687-696. DOI: 10.1197/jamia.M2470. <https://academic.oup.com/jamia/article/14/6/687/750453>

- [A28] Gal, M. S., & Rubinfeld, D. L. (2019). Data standardization. *New York University Law Review*, 94(4), 737-770. <https://www.nyulawreview.org/issues/volume-94-number-4/data-standardization/>
- [A29] Zheng, K., Ratwani, R. M., & Adler-Milstein, J. (2020). Studying workflow and workarounds in electronic health record-supported work to improve health system performance. *Annals of Internal Medicine*, 172(11 Suppl), S116-S122. DOI: 10.7326/M19-0871. <https://www.acpjournals.org/doi/10.7326/M19-0871>
- [A30] Bogaert, P., Verschuuren, M., Van Oyen, H., & Van Oers, H. (2021). Identifying common enablers and barriers in European health information systems. *Health Policy*, 125(12), 1517-1526. DOI: 10.1016/j.healthpol.2021.09.006. <https://www.sciencedirect.com/science/article/pii/S0168851021002396>

15.2 Industry Sources

- [I1] HIMSS Analytics. (2024). Analytics maturity assessment model (AMAM) global report. Healthcare Information and Management Systems Society. <https://www.himss.org/maturity-models/amam/>
- [I2] Snowdon, A. (2024). New analytics maturity adoption model pushes for digital transformation and data-driven decisions. *HIMSS*. <https://legacy.himss.org/news/new-analytics-maturity-adoption-model-pushes-digital-transformation-and-data-driven-decisions>
- [I3] Health Catalyst. (2020). The healthcare analytics adoption model: A roadmap to analytic maturity. <https://www.healthcatalyst.com/learn/insights/healthcare-analytics-adoption-model-roadmap-analytic-maturity>
- [I4] Berkshire Healthcare NHS Trust. (2024). Empowering citizen developers: Low-code success in healthcare. <https://ia.berkshirehealthcare.nhs.uk/citizen-developer-programme>
- [I5] Forrester Research. (2024). The total economic impact of Microsoft Power Apps. Forrester Consulting. <https://tei.forrester.com/go/microsoft/powerappstei/?lang=en-us>
- [I6] Oracle. (2024). The real cost of turnover in healthcare. <https://www.oracle.com/human-capital-management/cost-employee-turnover-healthcare/>
- [I7] Precedence Research. (2024). Healthcare analytics market size and forecast 2025 to 2034. <https://www.precedenceresearch.com/healthcare-analytics-market>
- [I8] Anthropic. (2025). Code modernization playbook: A practical guide to modernizing legacy systems with AI. <https://resources.anthropic.com/code-modernization-playbook>

[I9] IBM. (2022). Francisco Partners to Acquire IBM's Healthcare Data and Analytics Assets. *IBM Newsroom*. <https://newsroom.ibm.com/2022-01-21-Francisco-Partners-to-Acquire-IBMs-Healthcare-Data-and-Analytics-Assets>

[I10] LaVito, A. (2021). Haven, the Amazon-Berkshire-JPMorgan venture to disrupt healthcare, is disbanding after 3 years. *CNBC*. <https://www.cnbc.com/2021/01/04/haven-the-amazon-berkshire-jpmorgan-venture-to-disrupt-healthcare-is-disbanding-after-3-years.html>

[I11] American Health Information Management Association & NORC at the University of Chicago. (2023). Health information workforce survey report. <https://www.ahima.org/news-publications/press-room-press-releases/2023-press-releases/health-information-workforce-shortages-persist-as-ai-shows-promise-ahima-survey-reveals/>

16 Appendices

16.1 Appendix A: Healthcare Analytics Glossary

Term	Definition
AMAM	Analytics Maturity Assessment Model - HIMSS standard for measuring healthcare analytics capabilities
Clinical Terminology	Standardized vocabularies including ICD-10, CPT, SNOMED, and RxNorm used in healthcare data
Conversational AI	Artificial intelligence systems that enable natural language interaction for complex tasks
EHR	Electronic Health Record - digital version of patient medical records
HIMSS	Healthcare Information and Management Systems Society - healthcare IT standards organization
Institutional Memory	Collective organizational knowledge, expertise, and practices that enable effectiveness
NL2SQL	Natural Language to SQL - technology that converts spoken/written queries into database commands
Population Health	Analytics focused on health outcomes of groups of individuals rather than individual patients

Term	Definition
RAG	Retrieval Augmented Generation - AI approach combining information retrieval with text generation

16.2 Appendix B: HIMSS Analytics Maturity Assessment Model (AMAM) Stages

Stage	Name	Description	Key Capabilities
Stage 0	Data Collection	Basic data capture without integration	Manual data entry, paper records
Stage 1	Data Verification	Automated data validation and error checking	Basic quality controls, automated checks
Stage 2	Data Utilization	Standard reporting and basic analytics	Automated reports, dashboard creation
Stage 3	Automated Decision Support	Rule-based clinical and operational support	Clinical alerts, automated protocols
Stage 4	Population Health Analytics	Population-level analysis and intervention	Cohort identification, risk stratification
Stage 5	Predictive Analytics	Statistical modeling and predictive capabilities	Predictive models, trend analysis
Stage 6	Cognitive Analytics	Machine learning and AI integration	Advanced AI, pattern recognition
Stage 7	Precision Medicine	Personalized treatment and genomic integration	Personalized care, genomic analytics

16.3 Appendix C: Healthcare NL2SQL Query Examples

16.3.1 Patient Population Analysis

Natural Language: “Show me all diabetic patients over 65 who had an HbA1c test in the last 6 months”

Generated SQL:

```

1  SELECT p.patient_id, p.age, l.test_date, l.result_value
2  FROM patients p
3  JOIN lab_results l ON p.patient_id = l.patient_id
4  JOIN diagnoses d ON p.patient_id = d.patient_id
5  WHERE d.icd10_code LIKE 'E11%' -- Type 2 Diabetes
6    AND p.age > 65
7    AND l.test_name = 'HbA1c'
8    AND l.test_date >= DATE_SUB(CURRENT_DATE, INTERVAL 6 MONTH)

```

16.3.2 Operational Metrics

Natural Language: “What is the average length of stay for cardiac surgery patients this year?”

Generated SQL:

```

1  SELECT AVG(DATEDIFF(discharge_date, admission_date)) as avg_los
2  FROM admissions a
3  JOIN procedures p ON a.admission_id = p.admission_id
4  WHERE p.cpt_code IN ('33510', '33511', '33512') -- Cardiac surgery codes
5    AND a.admission_date >= '2025-01-01'
6    AND a.discharge_date IS NOT NULL

```

16.3.3 Quality Metrics

Natural Language: “How many patients were readmitted within 30 days of discharge for heart failure?”

Generated SQL:

```

1  SELECT COUNT(DISTINCT r.patient_id) as readmission_count
2  FROM (
3    SELECT a1.patient_id, a1.discharge_date, a2.admission_date
4    FROM admissions a1
5    JOIN admissions a2 ON a1.patient_id = a2.patient_id
6    JOIN diagnoses d ON a2.admission_id = d.admission_id
7    WHERE d.icd10_code LIKE 'I50%' -- Heart failure
8      AND a2.admission_date BETWEEN a1.discharge_date AND DATE_ADD(a1.
9        discharge_date, INTERVAL 30 DAY)
10     AND a1.admission_id != a2.admission_id
11   ) r

```

This work is licensed under a Creative Commons Attribution 4.0 International License.

Correspondence: <https://us.yuimedi.com/contact-us/> (*include “NL2SQL paper” in message*)